

LSTM Tagger

Philipp Windischhofer

January 5, 2017

The Setup

Goal

Train a binary neural-network based classifier that can distinguish between b - and non- b -jets, using the raw jet data as input.

- use *tracks* as the primary source of information
- number of tracks is unknown a-priori \rightarrow cannot use an architecture that expects a fixed number of inputs
- currently looking into recurrent neural networks / LSTM networks

Recurrent Neural Networks

The Workflow

Training

- match tracks to their associated jets (contained in different ROOT trees)
- for each track in the jet, feed all 8 available track parameters into the classifier network during training
 - ▶ use p_T ordering, i.e. hardest track first
- supervised training: provide a binary (0/1) output value for each jet (from MC)

Now running on Piz Daint:

- roughly $2\text{-}3 \times$ improvement in execution speed compared to PSI/Tier-3
- limited by Jet-Track-matching, which is handled by the CPU
- possible workaround: train multiple classifiers during the same run “in parallel”

The Workflow

Evaluation

- compare performance to cMVA tagger as “gold-standard”
- obtain ROC curves for both classifiers, correlation plots of the outputs
- currently: validation data is disjoint from training data, but from the same MC-run (i.e. contains a similar event signature)

Results so far

- trained a number of LSTM networks, scanned the hyperparameters:
 - ▶ number of nodes in each layer
 - ▶ number of layers
 - ▶ number of training epochs

Details of the training:

- read training data in chunks of 10k jets
- use 8k jets for training, 2k jets to monitor performance during each epoch

64 nodes / layer, 1 layer

- $\text{AUC}(\text{cMVA}) = \text{asdf}$
- $\text{AUC}(\text{RNN}) = \text{asdf}$

64 nodes / layer, 3 layers

- $\text{AUC}(\text{cMVA}) = \text{asdf}$
- $\text{AUC}(\text{RNN}) = \text{asdf}$

128 nodes / layer, 1 layer

- $\text{AUC}(\text{cMVA}) = \text{asdf}$
- $\text{AUC}(\text{RNN}) = \text{asdf}$

Future Experiments

-
- try removing track parameters to see how performance degrades
- more sophisticated preprocessing?
- different representation (other than the raw track data)?

Conclusions